

[Table of Contents](#) [Expand](#)

1 Problem Set 3: Neural Networks¶

This assignment requires a working IPython Notebook installation, which you should already have. If not, please refer to the instructions in the previous problem sets.

1.0.1 ** IMPORTANT NOTE FOR SUBMISSION **¶

1.0.2 In part 2 (programming) of this assignment, you DO NOT need to make modification to any existing code in this IPython Notebook. Instead you will implement your own simple neural network in the mlp.py file. To submit your answers, copy and paste the content of your mlp.py at the end of this IPython Notebook, then submit a single PDF version of this notebook.¶

Total: 100 points.


1.1 [30pts] Problem 1: Backprop in a simple MLP Table of Contents [Expand](#)

This problem asks you to derive all the steps of the backpropagation algorithm for a simple classification network. Consider a fully-connected neural network, also known as a multi-layer perceptron (MLP), with a single hidden layer and a one-node output layer. The hidden and output nodes use an elementwise sigmoid activation function and the loss layer uses cross-entropy loss:

$$f(z) = \frac{1}{1 + \exp(-z)}$$

$$L(\hat{y}, y) = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

The computation graph for an example network is shown below. Note that it has an equal number of nodes in the input and hidden layer (3 each), but, in general, they need not be equal. Also, to make the application of backprop easier, we show the *computation graph* which shows the dot product and activation functions as their own nodes, rather than the usual graph showing a single node for both.



The forward and backward computation are given below. NOTE: We assume no regularization, so you can omit the terms involving Ω .

The forward step is:

Require: Network depth, l
Require: $W^{(i)}, i \in \{1, \dots, l\}$, the weight matrices of the model
Require: $b^{(i)}, i \in \{1, \dots, l\}$, the bias parameters of the model
Require: x , the input to process
Require: y , the target output

```

 $h^{(0)} = x$ 
for  $k = 1, \dots, l$  do
   $a^{(k)} = b^{(k)} + W^{(k)} h^{(k-1)}$ 
   $h^{(k)} = f(a^{(k)})$ 
end for
 $\hat{y} = h^{(l)}$ 
 $J = L(\hat{y}, y) + \lambda \Omega(\theta)$ 

```

and the backward step is:

After the forward computation, compute the gradient on the output layer:

$$\mathbf{g} \leftarrow \nabla_{\hat{\mathbf{y}}} J = \nabla_{\hat{\mathbf{y}}} L(\hat{\mathbf{y}}, \mathbf{y})$$

for $k = l, l-1, \dots, 1$ do

Convert the gradient on the layer's output into a gradient on the pre-
nonlinearity activation (element-wise multiplication if f is element-wise):

$$\mathbf{g} \leftarrow \nabla_{\mathbf{a}^{(k)}} J = \mathbf{g} \odot f'(\mathbf{a}^{(k)})$$

Compute gradients on weights and biases (including the regularization term,
where needed):

$$\nabla_{\mathbf{b}^{(k)}} J = \mathbf{g} + \lambda \nabla_{\mathbf{b}^{(k)}} \Omega(\theta)$$

$$\nabla_{\mathbf{W}^{(k)}} J = \mathbf{g} \mathbf{h}^{(k-1)\top} + \lambda \nabla_{\mathbf{W}^{(k)}} \Omega(\theta)$$

Propagate the gradients w.r.t. the next lower-level hidden layer's activations:

$$\mathbf{g} \leftarrow \nabla_{\mathbf{h}^{(k-1)}} J = \mathbf{W}^{(k)\top} \mathbf{g}$$

end for

Table of Contents [Expand](#)

Write down each step of the backward pass explicitly for all layers, i.e. compute all intermediate gradients of J above, expressing them as a function of variables \mathbf{x} (a vector of inputs), y (a scalar label), $\mathbf{h}^{(k)}$ (layer k outputs), $\mathbf{W}^{(k)}$ (parameters in layer k), and $\mathbf{b}^{(k)}$ (bias parameters of layer k). Notice that $\mathbf{h}^{(k)}$ and $\mathbf{b}^{(k)}$ can be scalars or vectors depending on layer k , and $\mathbf{W}^{(k)}$ can be a matrix or a vector. NOTE: We will assume no regularization, so you can omit the terms involving Ω .

As an example, we will derive the first gradient for you, i.e. the gradient of the cross entropy loss w.r.t. its (scalar) input $\nabla_{\hat{\mathbf{y}}} J = \nabla_{\hat{\mathbf{y}}} L(\hat{\mathbf{y}}, y)$.

$$\nabla_{\hat{\mathbf{y}}} L(\hat{\mathbf{y}}, y) = \nabla_{\hat{\mathbf{y}}} [-y \ln(\hat{\mathbf{y}}) - (1 - y) \ln(1 - \hat{\mathbf{y}})] = \frac{\hat{\mathbf{y}} - y}{(1 - \hat{\mathbf{y}})\hat{\mathbf{y}}} = \frac{\mathbf{h}^{(2)} - y}{(1 - \mathbf{h}^{(2)})\mathbf{h}^{(2)}}$$

Please derive the remaining gradients, listed below.

Hint: you should substitute the updated values for the gradient \mathbf{g} in each step and simplify as much as possible.

Hint: Useful information about vectorized chain rule and backpropagation:

If you are struggling with computing the vectorized version of chain rule you may find this example helpful:

<https://web.stanford.edu/class/cs224n/readings/gradient-notes.pdf>

(<https://web.stanford.edu/class/cs224n/readings/gradient-notes.pdf>)

It also contains some helpful shortcuts for computing gradients.

[5pts] Q1.1: Derive $\nabla_{\mathbf{a}^{(2)}} J$, where $\mathbf{a}^{(2)}$ the (scalar) pre-nonlinearity activation of layer 2.

Hint: to get this scalar value, multiply the partial value \mathbf{g} , which should be equal to the $\nabla_{\hat{\mathbf{y}}} L(\hat{\mathbf{y}}, y)$ that we computed above, with the sigmoid derivative $f'(\mathbf{a}^{(2)})$, and note that $f(\mathbf{a}^{(2)}) = \mathbf{h}^{(2)}$.

Table of Contents [Expand](#)

$$\nabla_{a^{(2)}} J = \nabla_{a^{(2)}} L(\hat{y}, y) = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}}$$

Since the activation function is sigmoid $f(z) = \frac{1}{1+exp^{-z}}$,

$$f(z)' = \frac{1' \times (1+e^{-z}) - 1 \times (1+e^{-z})'}{(1+e^{-z})^2} = \frac{-1 \times e^{-z} \times -1}{(1+e^{-z})^2} = \frac{e^{-z}}{(1+e^{-z})^2} = \frac{1}{1+e^{-z}} \times \frac{e^{-z}}{1+e^{-z}} = \frac{1}{1+e^{-z}} \times \frac{1+e^{-z}-1}{1+e^{-z}} = f(z) \times (1-f(z))$$

I have $\hat{y} = h^{(2)} = f(a^{(2)})$, $\hat{y} = f(a^{(2)})$, when $f(x)$ is a sigmoid activation function. so.

$$\frac{\partial \hat{y}}{\partial a^{(2)}} = f(a^{(2)})' = f(a^{(2)}) \times (1 - f(a^{(2)})) = h^{(2)} \times (1 - h^{(2)}).$$

$$\text{So } \nabla_{a^{(2)}} J = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}} = \frac{h^{(2)} - y}{(1-h^{(2)})h^{(2)}} \times (h^{(2)} \times (1 - h^{(2)})) = h^{(2)} - y$$

[5pts] Q1.2: $\nabla_{b^{(2)}} J$

$$\nabla_{b^{(2)}} J = \nabla_{b^{(2)}} L(\hat{y}, y) = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial b^{(2)}}$$

$$a^{(2)} = h^{(1)} \times W^{(2)} + b^{(2)}, \text{ so } \frac{\partial a^{(2)}}{\partial b^{(2)}} = 1.$$

$$\text{Lastly, } \frac{\partial L(\hat{y}, y)}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial b^{(2)}} = (h^{(2)} - y) \times 1 = (h^{(2)} - y)$$

[5pts] Q1.3: $\nabla_{W^{(2)}} J$

Hint: this should be a vector, since $W^{(2)}$ is a vector.

$$\nabla_{W^{(2)}} J = \nabla_{W^{(2)}} L(\hat{y}, y) = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial W^{(2)}}$$

$$a^{(2)} = h^{(1)} \times W^{(2)} + b^{(2)}, \text{ so } \frac{\partial a^{(2)}}{\partial W^{(2)}} = h^{(1)}.$$

$$\text{Lastly, } \frac{\partial L(\hat{y}, y)}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial W^{(2)}} = (h^{(2)} - y) \times h^{(1)} = (h^{(2)} - y) \times h^{(1)}.$$

[5pts] Q1.4: $\nabla_{h^{(1)}} J$

$$\text{Note that } a^{(2)} = h^{(1)} \times W^{(2)} + b^{(2)}, \text{ so } \frac{\partial a^{(2)}}{\partial h^{(1)}} = W^{(2)}.$$

$$\nabla_{h^{(1)}} J = \nabla_{h^{(1)}} L(\hat{y}, y) = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial h^{(1)}}$$

$$\text{Lastly, } \frac{\partial L(\hat{y}, y)}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial h^{(1)}} = (h^{(2)} - y) \times W^{(2)} = (h^{(2)} - y) \times W^{(2)}.$$

[5pts] Q1.5: $\nabla_{b^{(1)}} J, \nabla_{W^{(1)}} J$

Table of Contents [Expand](#)

$$\nabla_{b^{(1)}} J = \nabla_{b^{(1)}} L(\hat{y}, y) = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial h^{(1)}} \times \frac{\partial h^{(1)}}{\partial a^{(1)}} \times \frac{a^{(1)}}{b^{(1)}}$$

$$\nabla_{W^{(1)}} J = \nabla_{W^{(1)}} L(\hat{y}, y) = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \times \frac{\partial \hat{y}}{\partial a^{(2)}} \times \frac{\partial a^{(2)}}{\partial h^{(1)}} \times \frac{\partial h^{(1)}}{\partial a^{(1)}} \times \frac{a^{(1)}}{W^{(1)}}$$

I have $h^{(1)} = f(a^{(1)})$, and $f(x)$ is also sigmoid function here, so $\frac{\partial h^{(1)}}{\partial a^{(1)}} = h^{(1)} \times (1 - h^{(1)})$.

From the equation, $a^{(1)} = W^{(1)} \times h^{(0)} + b^{(1)}$, $\frac{\partial a^{(1)}}{\partial W^{(1)}} = h^{(0)} = x$, $\frac{\partial a^{(1)}}{\partial b^{(1)}} = 1$.

Finally,

$$\frac{\partial L(\hat{y}, y)}{\partial h^{(1)}} \times \frac{\partial h^{(1)}}{\partial a^{(1)}} \times \frac{\partial a^{(1)}}{\partial W^{(1)}} = (h^{(2)} - y) \times W^{(2)} \times h^{(1)} \times (1 - h^{(1)}) \times h^{(0)}$$

$$\frac{\partial L(\hat{y}, y)}{\partial h^{(1)}} \times \frac{\partial h^{(1)}}{\partial a^{(1)}} \times \frac{\partial a^{(1)}}{\partial b^{(1)}} = (h^{(2)} - y) \times W^{(2)} \times h^{(1)} \times (1 - h^{(1)}) \times 1$$

[5pts] Q1.6 Briefly, explain how the computational speed of backpropagation would be affected if it did not include a forward pass

If backpropagation did not include a forward pass, its computational speed would be significantly slower due to redundant calculations. The forward pass pre-computes activation values and weighted sums, which are essential for efficiently computing gradients during backpropagation. Without these stored pre-calculated values, backpropagation would need to re-compute activations and intermediate values at each layer, leading to a significant increase in computational cost. This would not only slow down the training process but also require more memory access and operations, making neural network training inefficient. The inclusion of a forward pass allows backpropagation to reuse these computed values, improving both speed and efficiency.

1.2 [50pts] Problem 2 (Programming): Implementing a simple MLP¶

In this problem we will develop a neural network with fully-connected layers, or Multi-Layer Perceptron (MLP). We will use it in classification tasks.

In the current directory, you can find a file `m1p.py`, which contains the definition for class `TwoLayerMLP`. As the name suggests, it implements a 2-layer MLP, or MLP with 1 *hidden* layer. The hidden layer will have either the softmax or the RELU activation function. You will implement your code in the same file, and call the member functions in this notebook. Below is some initialization. The `autoreload` command makes sure that `m1p.py` is periodically reloaded.

In [1]:

```
# setup
import numpy as np
import matplotlib.pyplot as plt
from mlp import TwoLayerMLP

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Table of Contents [Expand](#)

Next we initialize a toy model and some toy data, the task is to classify five 4-d vectors.

In [2]:

```
# Create a small net and some toy data to check your implementations.
# Note that we set the random seed for repeatable experiments.
input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model(activ, std=1e-1):
    np.random.seed(0)
    return TwoLayerMLP(input_size, hidden_size, num_classes, std=std,
        activation=activ)

def init_toy_data():
    np.random.seed(1)
    X = 10 * np.random.randn(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

X, y = init_toy_data()
print('X = ', X)
print()
print('y = ', y)
```

```
X = [[ 16.24345364 -6.11756414 -5.28171752 -10.72968622]
 [ 8.65407629 -23.01538697 17.44811764 -7.61206901]
 [ 3.19039096 -2.49370375 14.62107937 -20.60140709]
 [-3.22417204 -3.84054355 11.33769442 -10.99891267]
 [-1.72428208 -8.77858418 0.42213747 5.82815214]]
```

```
y = [0 1 2 2 1]
```

1.2.1 [5pts] Q2.1 Forward pass: Sigmoid¶

Table of Contents [Expand](#)

The network takes a D-dimensional input vector and predicts C classes. It has the following architecture:

input - fully connected layer - ReLU (or sigmoid) - fully connected layer - softmax - Loss

In [3]:

```
# W1: First layer weights; has shape (D, H)
# b1: First layer biases; has shape (H,)
# W2: Second layer weights; has shape (H, C)
# b2: Second layer biases; has shape (C,)
```

Our 2-layer MLP uses a softmax output layer (**note**: this means that you don't need to apply a sigmoid on the output) and the multiclass cross-entropy loss to perform classification.

Softmax function:

For class j:

$$P(y_j|x) = \frac{\exp(z_j)}{\sum_{c=1}^C \exp(z_c)}$$

Where C is the number of classes and z is class-wise output of the network.

Multiclass cross-entropy loss function:

$$J = \frac{1}{m} \sum_{i=1}^m \sum_{c=1}^C [-y_c^{(i)} \log(P(y_c^{(i)}|x^{(i)}))]$$

$y_c^{(i)}$ is the label of i-th sample with respect to class c, and $y_c^{(i)} = 1$ for the ground truth class and 0 otherwise.

m is the number of inputs in a batch and C is the number of classes.

Please take a look at method `TwoLayerMLP.loss` in the file `mlp.py`. This function takes in the data and weight parameters, and computes the class scores (aka logits), the loss L , and the gradients on the parameters.

- Complete the implementation of forward pass (up to the computation of `scores`) for the sigmoid activation: $\sigma(x) = \frac{1}{1+\exp(-x)}$.

Note 1: Softmax cross entropy loss involves the [log-sum-exp operation](#)

(<https://en.wikipedia.org/wiki/LogSumExp>). This can result in numerical underflow/overflow. Read about the solution in the link, and try to understand the calculation of `loss` in the code.

Note 2: You're strongly encouraged to implement in a vectorized way and avoid using slower `for` loops. Note that most numpy functions support vector inputs.

Check the correctness of your forward pass below. The difference should be very small (<1e-6).

```
In [4]: net = init_toy_model('sigmoid')
loss, _ = net.loss(X, y, reg=0.1)
correct_loss = 1.182248
print(loss)
print('Difference between your loss and correct loss:')
print(np.sum(np.abs(loss - correct_loss)))
```

Table of Contents [Expand](#)

```
1.1822479803941373
Difference between your loss and correct loss:
1.9605862711102873e-08
```

1.2.2 [10pts] Q2.2 Backward pass: Sigmoid¶

- For sigmoid activation, complete the computation of `grads`, which stores the gradient of the loss with respect to the variables `W1`, `b1`, `W2`, and `b2`.

Notes: The outputs of the second fully-connected layer are the "scores" for each class. These are also often called "logits" and are the pre-nonlinearity values. After they go through the nonlinear softmax function they become probabilities. In the provided code they are stored in the N-by-C matrix "scores", i.e. for each data point there are C scores, one for each class.

In the backward pass you are asked to calculate the partial derivatives of the Loss with respect to each set of parameters and intermediate functions, using the backpropagation algorithm we covered in class. Each derivative is stored in a variable that starts with "d", so "dscores" refers to the derivative of the loss wrt the "scores" variable (the logits). Similarly, "dW2" stores the derivatives wrt the W2 parameters, etc.

Note that the derivatives should have the same shape as the variable, i.e. "dscores" should also be N-by-C.

"dW2" and "db2" have already been implemented for you.

In the first fully-connected layer, "dhidden" stores the derivatives wrt this layer's activations produced **after** the nonlinearity, i.e. "hidden" is the variable which is the input to the second fully-connected layer. This is the only hidden layer in this network, ie we do not observe its output during training.

Finally, when computing the rest of the backward pass, you need to include the local derivative of the nonlinear function of the first layer, which for this question is the sigmoid. Because the backward pass is modular, you only need to change this local derivative, the rest can remain the same!

By the way, mapping onto the first problem, "dscores" would be analogous to the derivative in Q1.1 and "dhidden" analogous to Q1.4, although the network structures are slightly different.

Now debug your backward pass using a numeric gradient check. Again, the differences should be very small.

In [7]:

```
# Use numeric gradient checking to check your implementation of the backward pass.
# If your implementation is correct, the difference between the numeric and analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.
from utils import eval_numerical_gradient

loss, grads = net.loss(X, y, reg=0.1)

# these should all be very small
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.1)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
    print('%s max relative error: %e'%(param_name, rel_error(param_grad_num, grads[param_name])))

W2 max relative error: 1.527005e-09
b2 max relative error: 5.553985e-11
W1 max relative error: 1.126756e-08
b1 max relative error: 2.035407e-06
```

Table of Contents

[Expand](#)

1.2.3 [5pts] Q2.3 Train the Sigmoid network¶

To train the network we will use stochastic gradient descent (SGD), implemented in `TwoLayerNet.train`. Then we train a two-layer network on toy data.

- Implement the prediction function `TwoLayerNet.predict`, which is called during training to keep track of training and validation accuracy.

You should get the final training loss around 0.1, which is good, but not too great for such a toy problem. One problem is that the gradient magnitude for `W1` (the first layer weights) stays small all the time, and the neural net doesn't get much "learning signals". This has to do with the saturation problem of the sigmoid activation function.

In [8]:

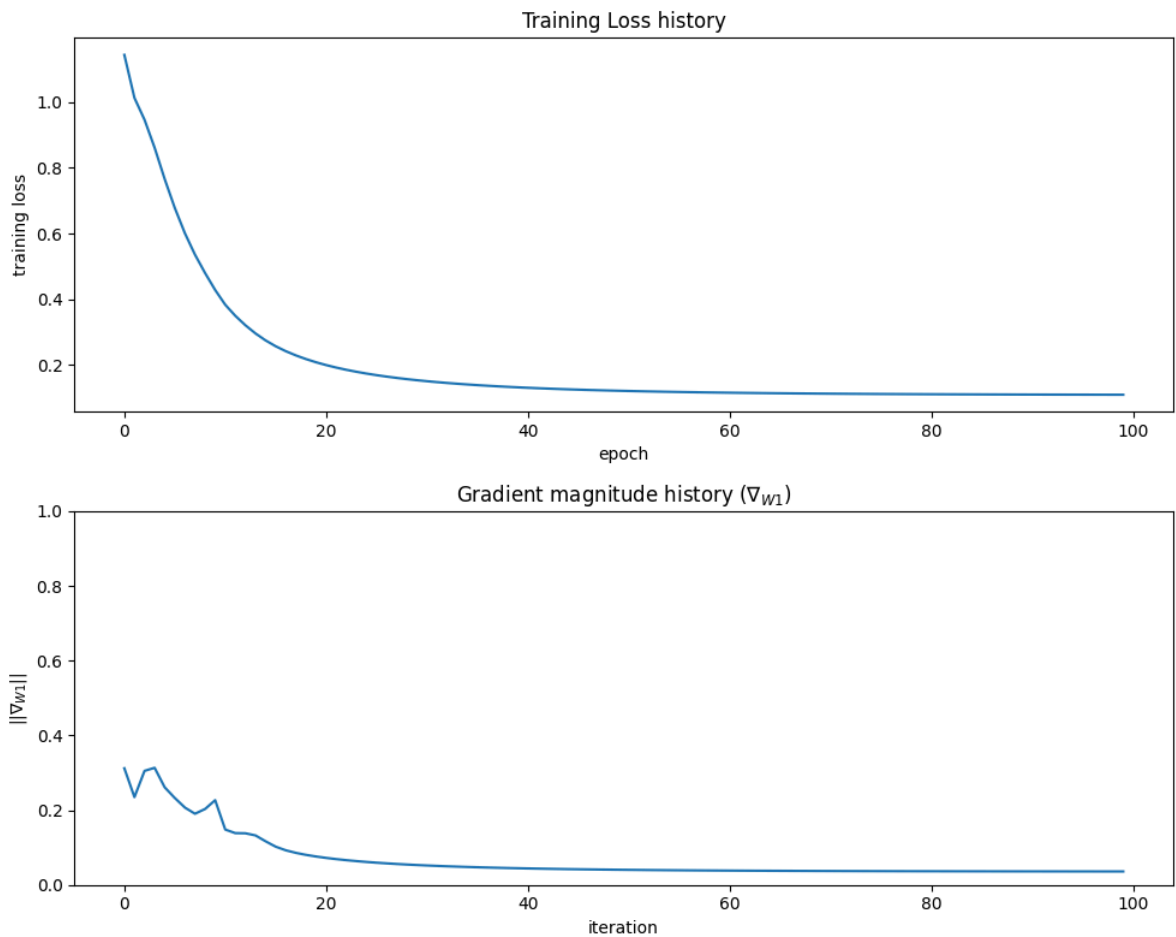
```

net = init_toy_model('sigmoid', std=1e-1)
stats = net.train(X, y, X, y,
                  learning_rate=0.5, reg=1e-5,
                  num_epochs=100, verbose=False)
print('Final training loss: ', stats['loss_history'][-1])

# plot the loss history and gradient magnitudes
fig, (ax1, ax2) = plt.subplots(2, 1)
ax1.plot(stats['loss_history'])
ax1.set_xlabel('epoch')
ax1.set_ylabel('training loss')
ax1.set_title('Training Loss history')
ax2.plot(stats['grad_magnitude_history'])
ax2.set_xlabel('iteration')
ax2.set_ylabel(r'$||\nabla_{W1}||$')
ax2.set_title('Gradient magnitude history ' + r'($\nabla_{W1}$)')
ax2.set_ylim(0,1)
fig.tight_layout()
plt.show()

```

Final training loss: 0.10926794610680679



1.2.4 [5pts] Q2.4 Using ReLU activation¶

Table of Contents [Expand](#)

The Rectified Linear Unit (ReLU) activation is also widely used: $ReLU(x) = \max(0, x)$.

- Complete the implementation for the ReLU activation (forward and backward) in `mlp.py`.
- Train the network with ReLU, and report your final training loss.

Make sure you first pass the numerical gradient check on toy data.

```
In [9]: net = init_toy_model('relu', std=1e-1)

loss, grads = net.loss(X, y, reg=0.1)
print('loss = ', loss) # correct_loss = 1.320973

# The differences should all be very small
print('checking gradients')
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.1)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
    print('%s max relative error: %e'%(param_name, rel_error(param_grad_num, grads[param_name])))

loss = 1.3037878913298202
checking gradients
W2 max relative error: 3.440708e-09
b2 max relative error: 4.447615e-11
W1 max relative error: 3.561318e-09
b1 max relative error: 2.738420e-09
```

Now that it's working, let's train the network. Does the net get stronger learning signals (i.e. gradients) this time? Report your final training loss.

In [10]:

```

net = init_toy_model('relu', std=1e-1)
stats = net.train(X, y, X, y,
                  learning_rate=0.1, reg=1e-5,
                  num_epochs=50, verbose=False)

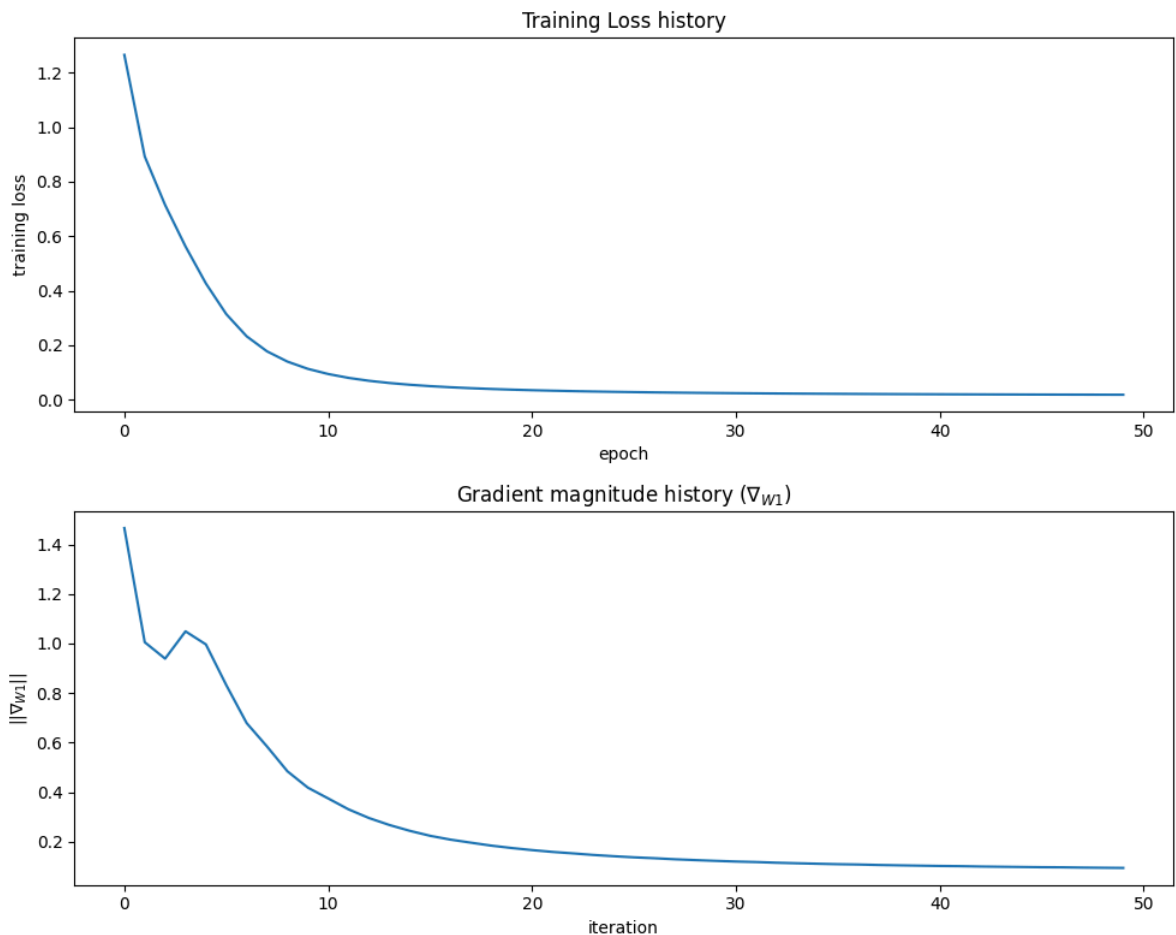
print('Final training loss: ', stats['loss_history'][-1])

# plot the loss history
fig, (ax1, ax2) = plt.subplots(2, 1)
ax1.plot(stats['loss_history'])
ax1.set_xlabel('epoch')
ax1.set_ylabel('training loss')
ax1.set_title('Training Loss history')
ax2.plot(stats['grad_magnitude_history'])
ax2.set_xlabel('iteration')
ax2.set_ylabel(r'$||\nabla_{W1}||$')
ax2.set_title('Gradient magnitude history ' + r'($\nabla_{W1}$)')
fig.tight_layout()
plt.show()

```

Table of Contents [Expand](#)

Final training loss: 0.0178562204869839



1.3 Load MNIST data¶

Table of Contents [Expand](#)

Now that you have implemented a two-layer network that works on toy data, let's try some real data. The MNIST dataset is a standard machine learning benchmark. It consists of 70,000 grayscale handwritten digit images, which we split into 50,000 training, 10,000 validation and 10,000 testing. The images are of size 28x28, which are flattened into 784-d vectors.

Note 1: the function `get_MNIST_data` requires the `scikit-learn` package. If you previously did anaconda installation to set up your Python environment, you should already have it. If you are using Google Colab, Colab should also have installed it for you. Otherwise, you can install it following the instructions here: <http://scikit-learn.org/stable/install.html> (<http://scikit-learn.org/stable/install.html>).

Note 2: If you encounter a `HTTP 500` error, that is likely temporary, just try again.

Note 3: Ensure that the downloaded MNIST file is 55.4MB (smaller file-sizes could indicate an incomplete download - which is possible)

```
In [11]: # load MNIST
from utils import get_MNIST_data
X_train, y_train, X_val, y_val, X_test, y_test = get_MNIST_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (50000, 784)
Train labels shape: (50000,)
Validation data shape: (10000, 784)
Validation labels shape: (10000,)
Test data shape: (10000, 784)
Test labels shape: (10000,)
```

1.3.1 Train a network on MNIST¶

Table of Contents [Expand](#)

We will now train a network on MNIST with 64 hidden units in the hidden layer. We train it using SGD, and decrease the learning rate with an exponential rate over time; this is achieved by multiplying the learning rate with a constant factor `learning_rate_decay` (which is less than 1) after each epoch. In effect, we are using a high learning rate initially, which is good for exploring the solution space, and using lower learning rates later to encourage convergence to a local minimum (or [saddle point](http://www.offconvex.org/2016/03/22/saddlepoints/) (<http://www.offconvex.org/2016/03/22/saddlepoints/>), which may happen more often).

- Train your MNIST network with 2 different activation functions: sigmoid and ReLU.

We first define some variables and utility functions. The `plot_stats` function plots the histories of gradient magnitude, training loss, and accuracies on the training and validation sets. The `show_net_weights` function visualizes the weights learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized. Both functions help you to diagnose the training process.

```
In [12]: input_size = 28 * 28
hidden_size = 64
num_classes = 10
```

Table of Contents [Expand](#)

Plot the loss function and train / validation accuracies

```
def plot_stats(stats):
    fig, (ax1, ax2, ax3) = plt.subplots(3, 1)
    ax1.plot(stats['grad_magnitude_history'])
    ax1.set_title('Gradient magnitude history ' + r'$(\nabla_{W1})$')
    ax1.set_xlabel('Iteration')
    ax1.set_ylabel(r'$||\nabla_{W1}||$')
    ax1.set_ylim(0, np.minimum(100, np.max(stats['grad_magnitude_history'])))

    ax2.plot(stats['loss_history'])
    ax2.set_title('Loss history')
    ax2.set_xlabel('Iteration')
    ax2.set_ylabel('Loss')
    ax2.set_ylim(0, 100)

    ax3.plot(stats['train_acc_history'], label='train')
    ax3.plot(stats['val_acc_history'], label='val')
    ax3.set_title('Classification accuracy history')
    ax3.set_xlabel('Epoch')
    ax3.set_ylabel('Classification accuracy')
    fig.tight_layout()
    plt.show()

# Visualize the weights of the network
from utils import visualize_grid
def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.T.reshape(-1, 28, 28)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()
```

1.3.2 [10pts] Q2.5 Train a Sigmoid network¶

Table of Contents [Expand](#)

```
In [13]: sigmoid_net = TwoLayerMLP(input_size, hidden_size, num_classes, activation='sigmoid', std=1e-1)

# Train the network
sigmoid_stats = sigmoid_net.train(X_train, y_train, X_val, y_val,
                                   num_epochs=20, batch_size=100,
                                   learning_rate=1e-3, learning_rate_decay=0.95,
                                   reg=0.5, verbose=True)

# Predict on the training set
train_acc = (sigmoid_net.predict(X_train) == y_train).mean()
print('Sigmoid final training accuracy: ', train_acc)

# Predict on the validation set
val_acc = (sigmoid_net.predict(X_val) == y_val).mean()
print('Sigmoid final validation accuracy: ', val_acc)

# Predict on the test set
test_acc = (sigmoid_net.predict(X_test) == y_test).mean()
print('Sigmoid test accuracy: ', test_acc)

# show stats and visualizations
plot_stats(sigmoid_stats)
show_net_weights(sigmoid_net)
```


/Users/jaylee/Documents/Documents - Jay's MacBook Pro/boston-univ/CAS CS 542/HW/pset3/mlp.py:86: RuntimeWarning: overflow encountered in exp
hidden = 1 / (1 + np.exp(-z1))

Table of Contents [Expand](#)

Epoch 1: loss 79.040004, train_acc 0.160000, val_acc 0.266300
 Epoch 2: loss 49.814996, train_acc 0.500000, val_acc 0.461100
 Epoch 3: loss 32.419904, train_acc 0.640000, val_acc 0.568300
 Epoch 4: loss 21.756599, train_acc 0.630000, val_acc 0.639700
 Epoch 5: loss 15.148895, train_acc 0.680000, val_acc 0.685000
 Epoch 6: loss 10.909900, train_acc 0.680000, val_acc 0.712600
 Epoch 7: loss 8.078106, train_acc 0.760000, val_acc 0.737900
 Epoch 8: loss 6.166522, train_acc 0.830000, val_acc 0.755600
 Epoch 9: loss 4.948016, train_acc 0.780000, val_acc 0.772900
 Epoch 10: loss 4.113118, train_acc 0.760000, val_acc 0.785000
 Epoch 11: loss 3.455138, train_acc 0.840000, val_acc 0.797000
 Epoch 12: loss 3.026239, train_acc 0.840000, val_acc 0.808100
 Epoch 13: loss 2.702231, train_acc 0.840000, val_acc 0.819600
 Epoch 14: loss 2.438965, train_acc 0.820000, val_acc 0.830900
 Epoch 15: loss 2.258613, train_acc 0.900000, val_acc 0.839900
 Epoch 16: loss 2.166625, train_acc 0.860000, val_acc 0.846800
 Epoch 17: loss 2.098843, train_acc 0.840000, val_acc 0.852800
 Epoch 18: loss 1.975990, train_acc 0.910000, val_acc 0.859300
 Epoch 19: loss 1.898398, train_acc 0.900000, val_acc 0.862300
 Epoch 20: loss 1.876564, train_acc 0.910000, val_acc 0.866400
 Sigmoid final training accuracy: 0.8721
 Sigmoid final validation accuracy: 0.8664
 Sigmoid test accuracy: 0.8639




Table of Contents						Expand

1.3.3 [10pts] Q2.6 Train a ReLU network¶

Table of Contents [Expand](#)

```
In [14]: relu_net = TwoLayerMLP(input_size, hidden_size, num_classes, activation='relu', std=1e-1)

# Train the network
relu_stats = relu_net.train(X_train, y_train, X_val, y_val,
                           num_epochs=20, batch_size=100,
                           learning_rate=1e-3, learning_rate_decay=0.95,
                           reg=0.5, verbose=True)

# Predict on the training set
train_acc = (relu_net.predict(X_train) == y_train).mean()
print('ReLU final training accuracy: ', train_acc)

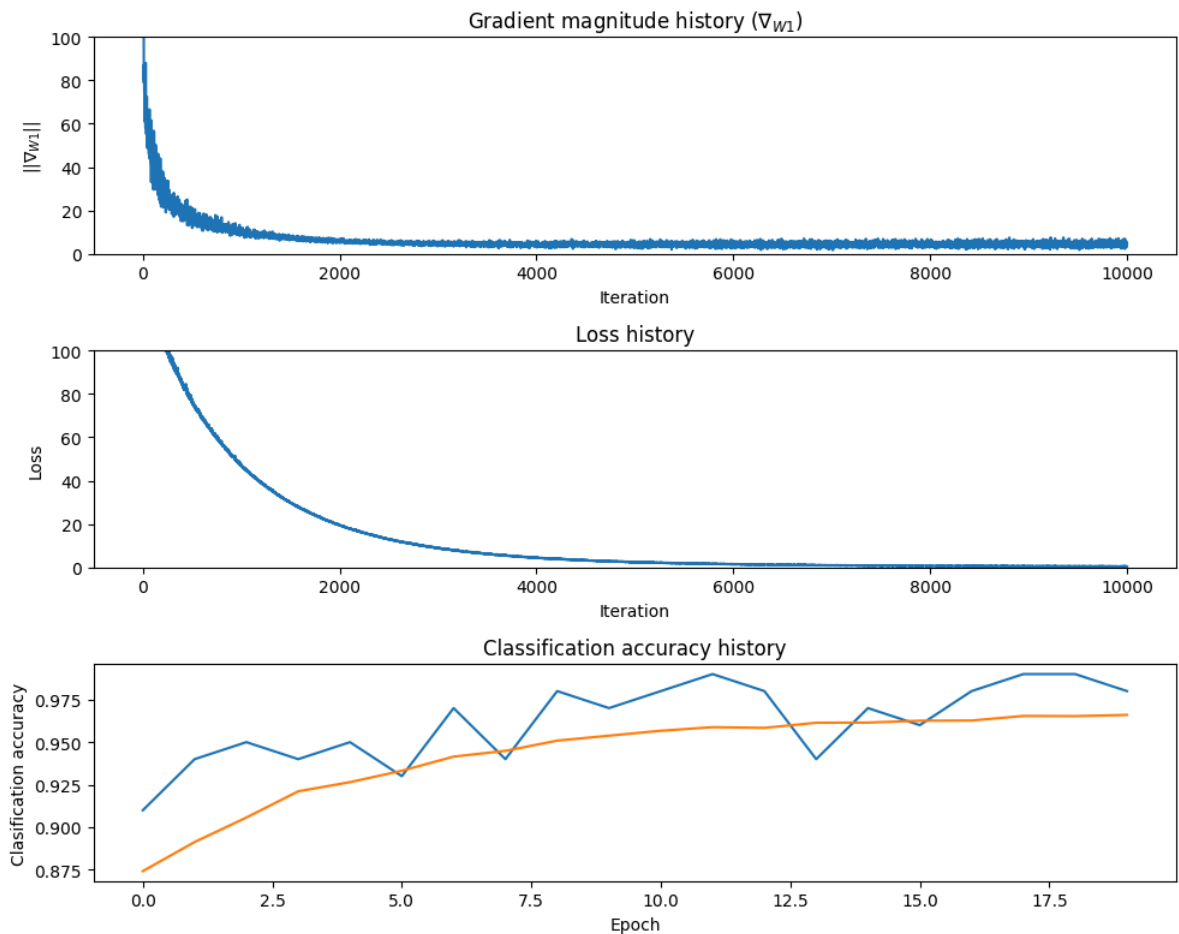
# Predict on the validation set
val_acc = (relu_net.predict(X_val) == y_val).mean()
print('ReLU final validation accuracy: ', val_acc)


# Predict on the test set
test_acc = (relu_net.predict(X_test) == y_test).mean()
print('ReLU test accuracy: ', test_acc)

# show stats and visualizations
plot_stats(relu_stats)
show_net_weights(relu_net)
```

Epoch 1: loss 77.144688, train_acc 0.910000, val_acc 0.874200
 Epoch 2: loss 47.540371, train_acc 0.940000, val_acc 0.891300
 Epoch 3: loss 29.869117, train_acc 0.950000, val_acc 0.905700
 Epoch 4: loss 19.522937, train_acc 0.940000, val_acc 0.911200
 Epoch 5: loss 13.036817, train_acc 0.950000, val_acc 0.926500
 Epoch 6: loss 8.990629, train_acc 0.930000, val_acc 0.933200
 Epoch 7: loss 6.142025, train_acc 0.970000, val_acc 0.941500
 Epoch 8: loss 4.526421, train_acc 0.940000, val_acc 0.944900
 Epoch 9: loss 3.237003, train_acc 0.980000, val_acc 0.950900
 Epoch 10: loss 2.411358, train_acc 0.970000, val_acc 0.953800
 Epoch 11: loss 1.849890, train_acc 0.980000, val_acc 0.956700
 Epoch 12: loss 1.393594, train_acc 0.990000, val_acc 0.958800
 Epoch 13: loss 1.162691, train_acc 0.980000, val_acc 0.958400
 Epoch 14: loss 0.994352, train_acc 0.940000, val_acc 0.961400
 Epoch 15: loss 0.815851, train_acc 0.970000, val_acc 0.961500
 Epoch 16: loss 0.682429, train_acc 0.960000, val_acc 0.962600
 Epoch 17: loss 0.625480, train_acc 0.980000, val_acc 0.962700
 Epoch 18: loss 0.488972, train_acc 0.990000, val_acc 0.965400
 Epoch 19: loss 0.451773, train_acc 0.990000, val_acc 0.965300
 Epoch 20: loss 0.404495, train_acc 0.980000, val_acc 0.966000
 ReLU final training accuracy: 0.9728
 ReLU final validation accuracy: 0.966
 ReLU test accuracy: 0.9658

Table of Contents

[Expand](#)




1.3.4 [5pts] Q2.71

Which activation function would you choose in practice? Why?

I would use the Relu activation function. First of all the data show the results it has a significant better results. Mnist dataset is the black and white data. By ignoring all the values below 0 could help find the black parts easier. I think this led to the better accuracy.

1.4 [20pts] Problem 3: Simple Regularization Methods [Table of Contents](#) [Expand](#)

You may have noticed the `reg` parameter in `TwoLayerMLP.loss`, controlling "regularization strength". In learning neural networks, aside from minimizing a loss function $\mathcal{L}(\theta)$ with respect to the network parameters θ , we usually explicitly or implicitly add some regularization term to reduce overfitting. A simple and popular regularization strategy is to penalize some *norm* of θ .

1.4.1 [10pts] Q3.1: L2 regularization

We can penalize the L2 norm of θ : we modify our objective function to be $\mathcal{L}(\theta) + \lambda \|\theta\|^2$ where λ is the weight of regularization. We will minimize this objective using gradient descent with step size η . Derive the update rule: at time $t + 1$, express the new parameters θ_{t+1} in terms of the old parameters θ_t , the gradient $g_t = \frac{\partial \mathcal{L}}{\partial \theta_t}$, η , and λ .

$$\lambda \|\theta\|^2 = \lambda \sum_i (\theta_i)^2$$

$$\frac{\partial \lambda \|\theta\|^2}{\partial \theta} = \frac{\partial \lambda \sum_i (\theta_i)^2}{\partial \theta} = \lambda \times \begin{bmatrix} \frac{\partial}{\partial \theta_0} \sum_i (\theta_i)^2 \\ \frac{\partial}{\partial \theta_1} \sum_i (\theta_i)^2 \\ \vdots \\ \frac{\partial}{\partial \theta_i} \sum_i (\theta_i)^2 \end{bmatrix} = \lambda \times \begin{bmatrix} 2\theta_0 \\ 2\theta_1 \\ \vdots \\ 2\theta_i \end{bmatrix}$$

Objective function is $\mathcal{L}(\theta) + \lambda \|\theta\|^2$, so the derivative of objective function respect to θ_t for each steps is $(\mathcal{L}(\theta) + \lambda \|\theta\|^2)' = \frac{\partial \mathcal{L}(\theta_t)}{\partial \theta_t} + \frac{\lambda \|\theta_t\|^2}{\partial \theta_t} = g_t + 2\lambda \theta_t$.

So the update rule at time $t + 1$ is

$$\theta_{t+1} = \theta_t - \eta \times (g_t + 2\lambda \theta_t)$$

1.4.2 [10pts] Q3.2: L1 regularization

Now let's consider L1 regularization: our objective in this case is $\mathcal{L}(\theta) + \lambda \|\theta\|_1$. Derive the update rule.

(Technically this becomes *Sub-Gradient* Descent since the L1 norm is not differentiable at 0. But practically it is usually not an issue.)

$$\lambda \|\theta\|_1 = \lambda \sum_i |\theta_i|$$

Table of Contents [Expand](#)

$$\frac{\partial}{\partial \theta_i} \lambda |\theta_i| = \begin{cases} \lambda, & \text{if } \theta_i > 0 \\ -\lambda, & \text{if } \theta_i < 0 \\ \text{Not differentiable,} & \text{if } \theta_i = 0 \end{cases}$$

Objective function is $\mathcal{L}(\theta) + \lambda \|\theta\|_1$, so the derivative of objective function respect to θ_t for each steps is $(\mathcal{L}(\theta) + \lambda \|\theta\|_1)' = \frac{\partial \mathcal{L}(\theta_t)}{\partial \theta_t} + \frac{\lambda \|\theta_t\|^2}{\partial \theta_t} = g_t + 2\lambda \theta$.

So the update rule at time $t + 1$ is

$$\theta_{t+1} = \theta_t - \eta \times (g_t + \lambda \times \text{sign}(\theta_t))$$

2 Supplementary Material¶

In []: *### copy and paste the content of your mlp.py here ###*

```
import numpy as np
import matplotlib.pyplot as plt
```

Table of Contents [Expand](#)

```
class TwoLayerMLP(object):
```

```
    """
```

A two-layer fully-connected neural network. The net has an input dimension of

N, a hidden layer dimension of H, and performs classification over C classes.

We train the network with a softmax loss function and L2 regularization on the

weight matrices. The network uses a ReLU nonlinearity after the first fully

connected layer.

In other words, the network has the following architecture:

input - fully connected layer - ReLU - fully connected layer - softmax

The outputs of the second fully-connected layer are the scores for each class.

```
    """
```

```
    def __init__(self, input_size, hidden_size, output_size, std=1e-4, activation='relu'):
```

```
        """
```

Initialize the model. Weights are initialized to small random values and

biases are initialized to zero. Weights and biases are stored in the

variable self.params, which is a dictionary with the following keys:

W1: First layer weights; has shape (D, H)

b1: First layer biases; has shape (H,)

W2: Second layer weights; has shape (H, C)

b2: Second layer biases; has shape (C,)

Inputs:

- input_size: The dimension D of the input data.

- hidden_size: The number of neurons H in the hidden layer.

- output_size: The number of classes C.

```
        """
```

```
        self.params = {}
```

```
        self.params['W1'] = std * np.random.randn(input_size, hidden_size)
```

```
        self.params['b1'] = np.zeros(hidden_size)
```

```
        self.params['W2'] = std * np.random.randn(hidden_size, output_size)
```

```
    e)
```

```
        self.params['b2'] = np.zeros(output_size)
```

```
        self.activation = activation
```

```
    def loss(self, X, y=None, reg=0.0):
```

```
        """
```

Compute the loss and gradients for a two layer fully connected neu


```

    ral
    network.

```

Inputs:

- *X*: Input data of shape (N, D) . Each $X[i]$ is a training sample.
- *y*: Vector of training labels. $y[i]$ is the label for $X[i]$, and each $y[i]$ is an integer in the range $0 \leq y[i] < C$. This parameter is optional; if it is not passed then we only return scores, and if it is passed then we instead return the loss and gradients.
- *reg*: Regularization strength.

Returns:

If *y* is None, return a matrix scores of shape (N, C) where scores $[i, c]$ is the score for class *c* on input $X[i]$.

If *y* is not None, instead return a tuple of:

- *loss*: Loss (data loss and regularization loss) for this batch of training samples.
- *grads*: Dictionary mapping parameter names to gradients of those parameters with respect to the loss function; has the same keys as *self.params*.

```

    """
    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    _, C = W2.shape
    N, D = X.shape

    # Compute the forward pass
    #####
    #####
    # write your own code where you see [PLEASE IMPLEMENT]
    #
    # Perform the forward pass, computing the class scores for the input.
    # Store the result in the scores variable, which should be an array of
    # shape (N, C).
    #####
    #####
    z1 = np.dot(X, W1) + b1 # 1st layer activation, N*H

    # 1st layer nonlinearity, N*H
    if self.activation == 'relu':
        hidden = np.maximum(0, z1)
    elif self.activation == 'sigmoid':
        hidden = 1 / (1 + np.exp(-z1))
    else:
        raise ValueError('Unknown activation type')

    # [PLEASE IMPLEMENT] 2nd layer activation, N*C

```

Table of Contents [Expand](#)

`# hint: involves W2, b2`
`z2 = np.dot(hidden, W2) + b2`

`# No need to make Softmax function, I'll make softmax to figure it out....`
`# a2 = np.exp(z2) / np.sum(np.exp(z2), axis = 1, keepdims=True)`

`scores = z2 # Scores N*C`
`#####`
`#####`
`#`
`##### END OF YOUR CODE`
`#####`
`#####`

`# If the targets are not given then jump out, we're done`
`if y is None:`
 `return scores`

`# cross-entropy loss with log-sum-exp`
`A = np.max(scores, axis=1) # N*1 All the predicted classes`
`F = np.exp(scores - A.reshape(N, 1)) # N*C the predicted one will`
`be 0 and rests will be negative values`
`P = F / np.sum(F, axis=1).reshape(N, 1) # N*C changed into positive values`
`loss = np.mean(-np.choose(y, scores.T) + np.log(np.sum(F, axis=1))`
`+ A)`
`# add regularization terms`
`loss += 0.5 * reg * np.sum(W1 * W1)`
`loss += 0.5 * reg * np.sum(W2 * W2)`

`# Backward pass: compute gradients`
`grads = {}`
`#####`
`#####`
`# write your own code where you see [PLEASE IMPLEMENT]`
`#`
`# Compute the backward pass, computing the derivatives of the weights`
`and biases. Store the results in the grads dictionary. For example,`
`# grads['W1'] should store the gradient on W1, and be a matrix of`
`same size.`
`# You should define the hidden variable in the part where you implement the`
`different activation functions. "hidden" is the output after you apply`
`the activation function for the hidden layer.`
`# Hint: you should apply different activation functions on "z1" and get "hidden".`
`#####`
`#####`

`# output layer`
`dscore = (P - np.eye(C)[y]) # partial derivative of loss wrt. the logits (dL/dz)`
`dW2 = np.dot(hidden.T, dscore) / N # partial derivative of loss w`

Table of Contents [Expand](#)

```

rt. W2
    db2 = np.mean(dscore, axis=0)      # partial derivation of loss wr
t. b2

# hidden layer
dhidden = np.dot(dscore, W2.T)

if self.activation == 'relu':
    dz1 = dhidden * (z1 > 0).astype(float)
elif self.activation == 'sigmoid':
    dz1 = dhidden * hidden * (1 - hidden)
else:
    raise ValueError('Unknown activation type')

# first layer
dW1 = np.dot(X.T, dz1) / N
db1 = np.mean(dz1, axis=0)
#####
#####
#                                     END OF YOUR CODE
#####
#####

grads['W2'] = dW2 + reg * W2
grads['b2'] = db2
grads['W1'] = dW1 + reg * W1
grads['b1'] = db1
return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_epochs=10,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
        X[i] has label c, where 0 <= c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning rate
        after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """
    num_train = X.shape[0]
    iterations_per_epoch = max(int(num_train / batch_size), 1)
    epoch_num = 0

```

Table of Contents [Expand](#)

Table of Contents [Expand](#)

```

# Use SGD to optimize the parameters in self.model
loss_history = []
grad_magnitude_history = []
train_acc_history = []
val_acc_history = []

np.random.seed(1)
for epoch in range(num_epochs):
    # fixed permutation (within this epoch) of training data
    perm = np.random.permutation(num_train)

    # go through minibatches
    for it in range(iterations_per_epoch):
        X_batch = None
        y_batch = None

        # Create a random minibatch
        idx = perm[it*batch_size:(it+1)*batch_size]
        X_batch = X[idx, :]
        y_batch = y[idx]

        # Compute loss and gradients using the current minibatch
        loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
        loss_history.append(loss)

        # do gradient descent
        for param in self.params:
            self.params[param] -= grads[param] * learning_rate

        # record gradient magnitude (Frobenius) for W1
        grad_magnitude_history.append(np.linalg.norm(grads['W1']))

    # Every epoch, check train and val accuracy and decay learning
    rate.

    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)
    if verbose:
        print('Epoch %d: loss %f, train_acc %f, val_acc %f'%(
            epoch+1, loss, train_acc, val_acc))

    # Decay learning rate
    learning_rate *= learning_rate_decay

    return {
        'loss_history': loss_history,
        'grad_magnitude_history': grad_magnitude_history,
        'train_acc_history': train_acc_history,
        'val_acc_history': val_acc_history,
    }

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict label

```

s for
 data points. For each data point we predict scores for each of the
 C
 classes, and assign each data point to the class with the highest
 score.

Table of Contents [Expand](#)

Inputs:

– X: A numpy array of shape (N, D) giving N D-dimensional data points to
 classify.

Returns:

– y_pred: A numpy array of shape (N,) giving predicted labels for
 each of
 the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
 to have class c, where $0 \leq c < C$.

```
#####
#####
# hint: it should be very easy
scores = self.loss(X)
y_pred = np.argmax(scores, axis=1)
#####
#####

return y_pred
```